

ECRAG: Enhanced Corrective Retrieval Augmented Generation with Sub-Document Embedding And Efficient Knowledge Detection

Bob Bobga Nti

National College of Ireland

Abstract

The **Enhanced Corrective Retrieval Augmented Generation (ECRAG)** framework is introduced as an innovative approach to enhance the accuracy, relevance, and efficiency of information retrieval and response generation in knowledge augmented language models. Building upon existing systems like **Retrieval Augmented Generation (RAG)** and **Corrective Retrieval Augmented Generation (CRAG)**, ECRAG incorporates hierarchical sub-document embeddings, a dynamic web integration mechanism, and a refined knowledge detection process to improve retrieval quality and minimize hallucinations.

ECRAG was evaluated using a domain-specific knowledge base comprising 250 arXiv AI research papers, with performance benchmarks against RAG and CRAG. Results demonstrated ECRAG's superior retrieval accuracy (up to 99%), significantly reduced response hallucination rates (5%), and enhanced response coherence and quality (rated 4.6 out of 5 by human evaluators). The Knowledge Refinement Mechanism was pivotal in ensuring generated content was factually correct, contextually relevant, and aligned with user queries. Additionally, ECRAG's capability to perform dynamic web searches when internal content was insufficient augmented its versatility in handling complex and emerging topics.

Despite its advancements, ECRAG has limitations, including dependency on the quality of external sources and computational overhead from integrating multiple components. Future work aims to enhance web content verification, improve scalability through distributed embedding storage, and incorporate user feedback for further model refinement.

ECRAG represents a significant advancement in retrieval-augmented generation, leveraging corrective retrieval mechanisms and dynamic external augmentation to create robust, adaptable, and context-aware language models suitable for applications such as research assistance, specialized customer support, and educational tools.

Keywords: Retrieval-Augmented Generation (RAG), Sub-Document Embeddings, Knowledge Detection, Hallucination Mitigation, Information Retrieval

1 Introduction

The rapid advancement of Artificial Intelligence (AI), particularly in Natural Language Processing (NLP), has led to the widespread adoption of Large Language Models (LLMs) for tasks like content generation, summarisation, and automated question answering. Despite their capabilities, LLMs often produce responses lacking factual grounding, a phenomenon known as hallucination (Ji et al., 2023). Hallucinations undermine the reliability and utility of LLMs, especially in applications demanding high factual accuracy.

To mitigate hallucinations, Retrieval Augmented Generation (**RAG**) systems have been developed, integrating an information retrieval component to anchor generative outputs in relevant source documents (Lewis et al., 2020). RAG retrieves documents or passages based on a query and feeds the retrieved context to the generative model, thereby reducing hallucinations by grounding responses in external knowledge. However, RAG systems face challenges with retrieval speed and accuracy, particularly with large or complex datasets, leading to degraded response quality (Borgeaud et al., 2022).

In attempting to address the performance quality, the Corrective Retrieval Augmented Generation (**CRAG**) framework was proposed by Yan et al. (2023). CRAG introduces a corrective mechanism that re-ranks retrieved passages based on relevance before passing them to the generative model, enhancing response grounding and coherence. Despite these improvements, CRAG still encounters limitations in efficiently detecting contextually relevant knowledge and ensuring high retrieval accuracy at scale.

This research presents Enhanced Corrective Retrieval Augmented Generation (**ECRAG**), an extension of RAG and CRAG and designed to overcome their limitations through hierarchical sub-document embeddings and integrating a robust Knowledge Detection Mechanism (**KDM**). ECRAG segments documents into smaller, semantically coherent chunks and subchunks thereby allowing for more precise query-aligned retrieval. By employing a dual-layer scoring mechanism combining **Cosine Similarity (CS)** and **Best Matching 25 (BM25)** based keyword matching and ranking, ECRAG filters irrelevant or weakly aligned data early in the retrieval process, reducing noise and enhancing response quality. These approaches significantly improves both retrieval speed and accuracy, offering a scalable solution for large-scale retrieval tasks.

2 Related Work

2.1 Retrieval Augmented Generation (RAG)

Retrieval-Augmented Generation (RAG) systems combine information retrieval with language model generation to enhance factual accuracy in outputs. The framework retrieves relevant context from a document corpus and integrates this context into the generative process. Lewis et al. (2020) demonstrated that RAG systems reduce hallucinations and improve response quality by grounding the generation process in actual documents. However, the retrieval component in standard RAG systems struggles with speed and scalability as the size of the corpus increases. Moreover, the reliance on fixed chunking strategies can result in the omission of critical information buried within large text fragments.

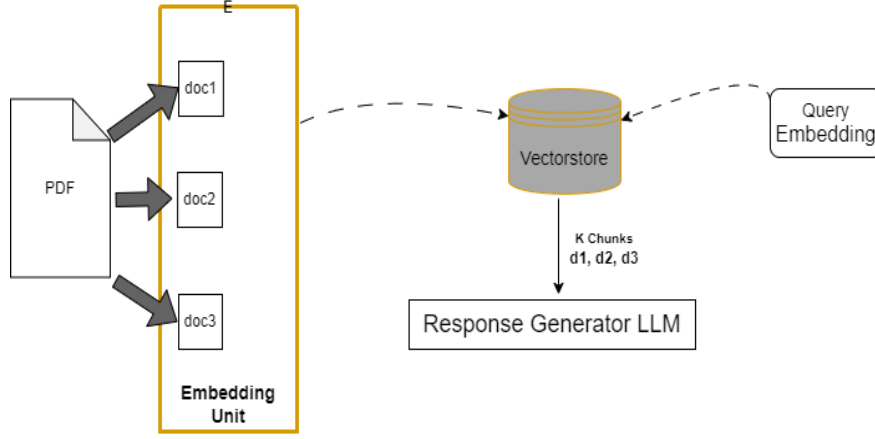


Fig. 1: Standard RAG Architecture

2.2 Corrective Retrieval - Augmented Generation (CRAG)

To address the limitations of RAG, Yan et al. (2023) proposed the Corrective Retrieval-Augmented Generation (CRAG) framework. CRAG introduces a corrective re-ranking mechanism that evaluates retrieved passages for relevance before feeding them to the generative model. This additional step improves grounding and coherence in responses. However, CRAG inherits certain limitations from RAG, particularly in handling ambiguous queries and ensuring high retrieval accuracy when processing large or unstructured datasets.

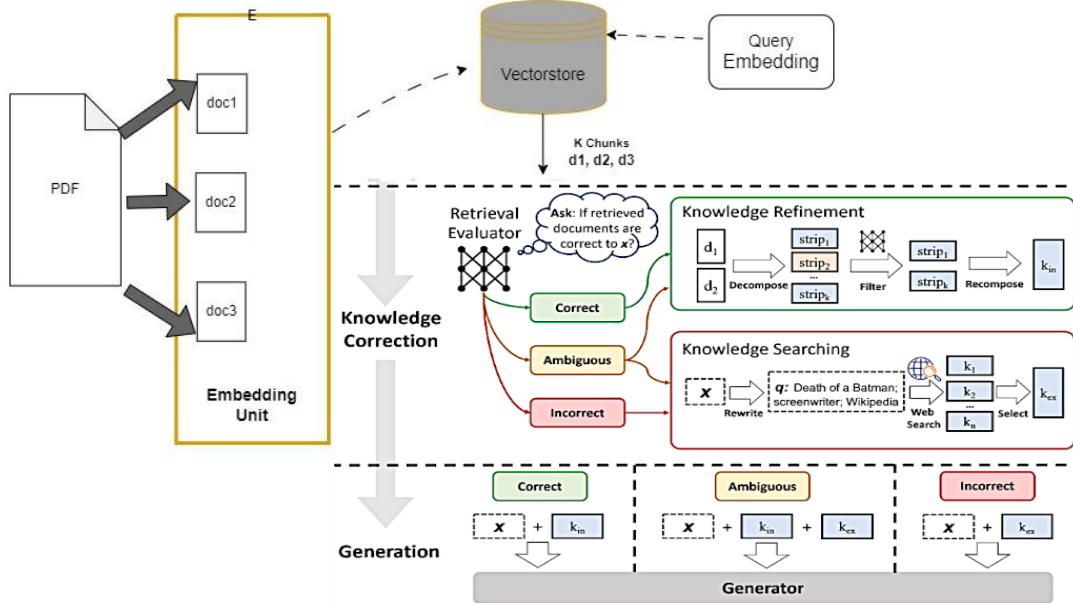


Figure 2: RAG with CRAG.

2.3 Document Chunking and Sub-Chunking

Effective retrieval in RAG-based systems often depends on how documents are chunked. Most systems use fixed-length token or sentence-based chunking, which may fail to capture fine-grained details relevant to the query (Borgeaud et al., 2022). Recent advancements advocate for hierarchical document embeddings, where documents are divided into chunks and further into subchunks, enabling more granular retrieval.

2.4 Knowledge Detection and Refinement

Knowledge detection mechanisms help identify whether a query can be answered using internal domain knowledge or requires external augmentation. Traditional approaches rely on LLM confidence scores or simple keyword matching, but these methods are prone to false positives and negatives. Recent research combines lexical (e.g., BM25) and semantic (e.g., cosine similarity) approaches for more robust classification.

3 Research Methodology

3.1 Overview

The aim of this research is to address key challenges in retrieval augmented generation systems. In particular, the study investigates how hierarchical sub-document embedding and a dynamic Knowledge Detection Mechanism (KDM) can improve retrieval accuracy, reduce hallucinations, and maintain scalability. The research methodology is structured into three main phases:

- **System Development:** Establish the theoretical underpinnings and design the overall retrieval framework.
- **Experimental Evaluation:** Empirically test the system against baseline models (RAG and CRAG) using domain-specific datasets.
- **Comparative Analysis:** Analysing performance metrics to understand the improvements contributed by each component of the system.

3.2 Research Questions and Objectives

The primary research questions include:

- How does the ECRAG framework improve retrieval accuracy and reduce hallucination compared to existing models?
- What is the impact of hierarchical embeddings on scalability and query latency?
- How effectively does the KDM balance internal and external knowledge integration?

Objectives are set to quantitatively assess the performance using standard retrieval metrics and human evaluation scores.

3.3 Data Collection and Preprocessing

Data is collected from a curated corpus of 250 domain-specific AI research papers.

Preprocessing steps include:

- Extraction of text from PDFs.
- Clean-up of formatting and removal of irrelevant sections (e.g. References).
- Segmentation into chunks and further into subchunks using sentence-aware algorithms.

Mathematical formulations are utilised to define the chunking process, ensuring that each chunk maintains semantic integrity. (See Section 3.4 for related equations.)

3.4 Evaluation Metrics

The evaluation is based on the following standard metrics:

- **Precision@k (P):**
 $P = (\text{Number of Relevant Documents Retrieved}) / (\text{Total Number of Retrieved Documents})$
- **Recall@k (R):**
 $R = (\text{Number of Relevant Documents Retrieved}) / (\text{Total Number of Relevant Documents in the Dataset})$
- **F1-Score:**
 $F1 = 2 \cdot (P \cdot R) / (P + R)$

Additionally, the **Hallucination Rate (HR)** is defined as:

$$HR = (\text{Number of Responses with Hallucinations}) / (\text{Total Number of Responses})$$

Scalability is measured by analysing query latency,

which empirically scales as $T \propto O(n \log n)$ with respect to the number of documents, n .

4 Design Specification

This section details the conceptual design and architectural blueprint of the ECRAG framework. It focuses on the system's modular structure, describing each component and its interrelation within the overall architecture.

4.1 System Overview

ECRAG is built on a modular, hierarchical architecture designed to overcome limitations of conventional retrieval systems. The system is designed to:

- Segment documents into manageable chunks and subchunks for precise retrieval.
- Employ a dual-layer retrieval pipeline that integrates semantic and keyword-based matching.
- Dynamically decide whether to utilise internal or external sources based on query ambiguity via the Knowledge Detection Mechanism (KDM) of the Knowledge Detection Unit (KDU).

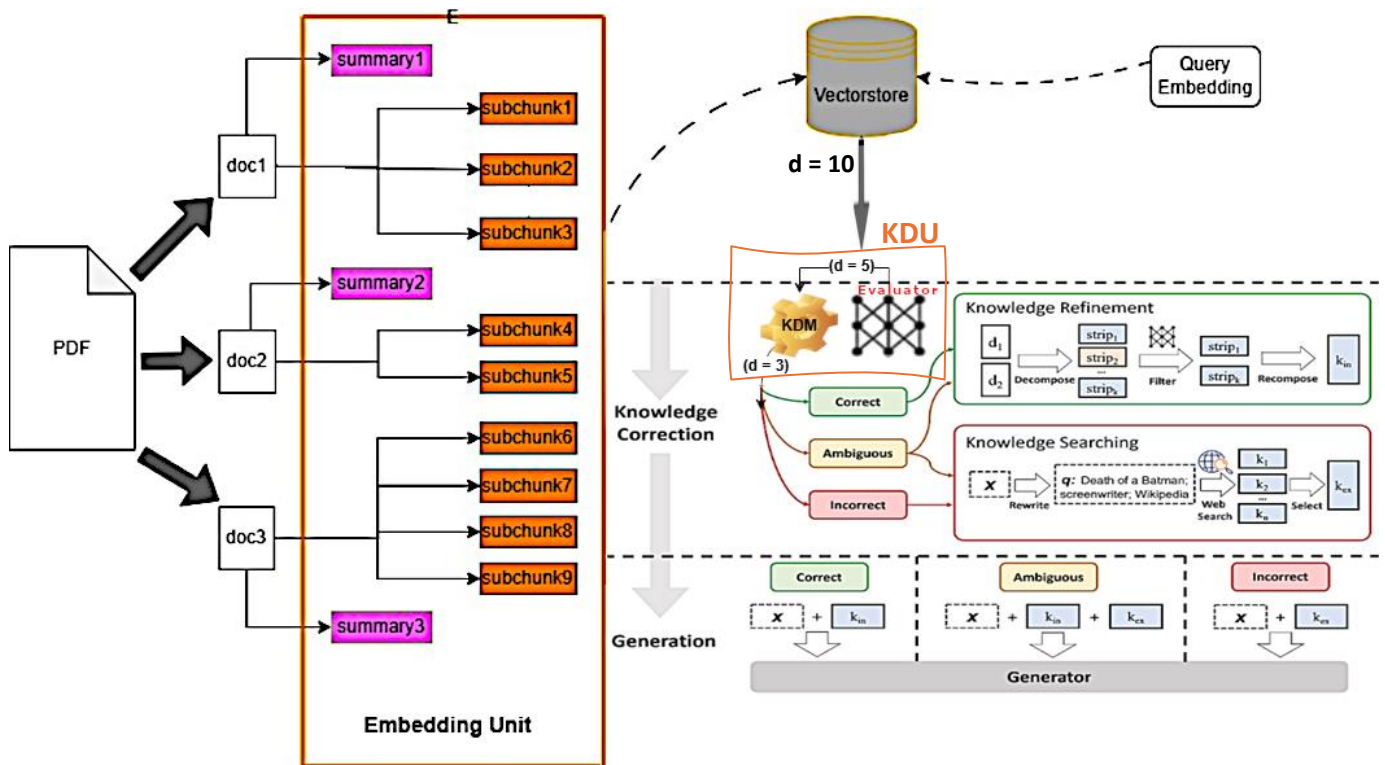


Fig. 3: ECRAAG Architecture

4.2 Architecture Components

The architecture comprises the following key modules:

I. Data Processing Module:

Ingests documents, cleans the text, and segments it into chunks and subchunks.

To ensure optimal segmentation of the document for efficient retrieval, the following mathematical formulations are employed in our chunking and subchunking process.

a. Chunking

Let T be the total number of characters in the document (pdf file). The following parameters are defined:

- Maximum chunk size: L (e.g. 2500 characters)
- Minimum chunk size: M (e.g. 300 characters)
- Chunk overlap: O (e.g. 100 characters)

The approximate number of chunks, N , is estimated as:

$$N \approx [(T - L) / (L - O)] + 1$$

Note: If a chunk's size falls below M , it is merged with an adjacent chunk, thereby adjusting the effective chunk size.

b. Subchunking Within Each Chunk

For each chunk of length L , the segmentation into subchunks is defined by:

- Subchunk size: l (e.g. 300 characters)
- Subchunk overlap: o (e.g. 50 characters)

The approximate number of subchunks, n , per chunk is given by:

$$n \approx \lceil (L - l) / (l - o) \rceil + 1$$

This formula guarantees that each subchunk maintains semantic continuity with an overlap of o characters to preserve context.

c. Summarisation Parameters for Chunks

Each chunk is summarised using a dedicated summarisation pipeline. The summary length is determined as a fraction of the chunk's length. Let $|\text{chunk}|$ denote the number of characters in the chunk. The token-based summary length parameters are defined as:

- Maximum summary length (S_{max}): $S_{\text{max}} = 0.50 \times |\text{chunk}| / 4$
- Minimum summary length (S_{min}): $S_{\text{min}} = 0.20 \times |\text{chunk}| / 4$

Here, the division by 4 is an approximate conversion from characters to tokens.

d. Counters and Identifiers

The system utilises global counters to ensure unique identification of each document, chunk, and subchunk:

- Global document counter: increments per document processed.
- Global chunk counter: increments for each chunk across all documents.
- Global subchunk counter: increments for each subchunk across all chunks.

These counters (i.e. `global_document_counter`, `global_chunk_counter`, and `global_subchunk_counter`) ensure traceability and uniqueness within the hierarchical structure.

II. Embedding and Indexing Module:

Generates embeddings for both chunk summaries and subchunks using state-of-the-art models and indexes them using FAISS.

III. Retrieval Module:

Implements a hierarchical retrieval strategy, initially retrieving chunk summaries and then refining results with subchunk data.

- ❖ **Coarse Retrieval:** FAISS (IndexFlatL2) retrieves the top k_1 chunk summaries based on cosine similarity.
- ❖ **Fine-Grained Search:** A new FAISS(IndexFlatL2) of all the subchunk embeddings of all the subchunks from the retrieved chunks. The top k_2 subchunks are then retrieved.

IV. Knowledge Detection Module (KDM):

A pivotal element of ECRAg's knowledge detection mechanism. This is the knowledge algorithm which computes a hybrid confidence score (C) that integrates

both semantic similarity and keyword-based matching. The confidence score is computed as:

$$C = 0.7 \cdot CS + 0.3 \cdot BM25$$

Where:

✓ **Cosine Similarity (CS):**

$$CS = (A \cdot B) / (\|A\| \cdot \|B\|)$$

Here, A and B are the embedding vectors of the query and the document respectively, and $\|A\|$ denotes the Euclidean norm of vector A.

✓ **BM25 Score:**

$$BM25 = \sum [IDF(q_i) \cdot ((f(q_i, D) \cdot (k_1 + 1)) / (f(q_i, D) + k_1 \cdot (1 - b + b \cdot (|D|/avgD))))]$$

where:

- ✓ q_i is the i-th term in the query,
- ✓ $f(q_i, D)$ is the term frequency of q_i in document D,
- ✓ $|D|$ is the length of document D,
- ✓ avgD is the average document length in the corpus,
- ✓ k_1 and b are tunable parameters (typically $k_1 = 1.2$ and $b = 0.75$),
- ✓ $IDF(q_i)$ represents the inverse document frequency of term q_i .

V. **Generative Response Module:**

Integrates retrieved subchunks and any additional external content into a coherent prompt for the language model.

4.3 Data Flow and Information Processing

The data flow is carefully designed to ensure minimal overlap between components. The process begins with document ingestion, followed by segmentation, embedding, and indexing. During query processing, the KDM directs the system to the appropriate retrieval pathway, ensuring that only high relevance subchunks are passed on to the generative module. This conceptual separation ensures clarity in design while allowing for optimised performance.

5 Implementation

The Implementation section provides a detailed account of the technical realisation of the ECRAAG framework. It covers the programming tools, libraries, and pseudocode utilised, with emphasis on reproducibility and practical execution.

5.1 Technical Environment

The system is implemented in Python (v3.9+) utilising:

- **Libraries:**
 - ✓ transformers (for embeddings and summarisation),

- ✓ FAISS (for efficient nearest neighbour search),
- ✓ PyTorch (for deep learning operations),
- NLTK (for sentence tokenisation), and
- ✓ Additional libraries such as re, uuid, and pickle for text processing and object serialisation.
- **Hardware:**

GPU acceleration is utilised where available; otherwise, the system defaults to CPU processing.

5.2 Core Implementation Details

The practical implementation is divided into the following components:

- **Document Processing and Chunking:**

Code is structured to open PDF files, clean the text, and segment it into chunks.

 - ✓ Each chunk is summarised using a pre-trained summarisation model (facebook/bart-large-cnn).
 - ✓ Mathematical formulations for determining the number of chunks are implemented as described in Section 4.2.
- **Subchunking Process:**

Each chunk is further divided into subchunks using sentence-level tokenisation.

Global counters (global_chunk_counter, global_subchunk_counter) ensure unique identification of each entity.

The pseudocode (provided in the Appendix) outlines the logic, including merging small chunks and creating metadata for traceability.
- **Embedding and Indexing:**

Embeddings for both chunks and subchunks are computed using a dedicated embedding model.

These embeddings are then indexed using FAISS for efficient retrieval.
- **Hybrid Retrieval and Confidence Scoring:**

During query processing, the system computes a hybrid confidence score, integrating cosine similarity and BM25, as specified earlier: $C = 0.7 \cdot CS + 0.3 \cdot BM25$.

The implementation of this scoring function is key to ensuring that the most relevant subchunks are selected for generating responses.

5.3 Pseudocode and Code Organisation

Detailed pseudocode is provided for each major process:

- **Chunking and Subchunking Algorithms:**

Outlines the segmentation process including merging logic for small chunks.

I. Chunking

```
FUNCTION _create_chunks(text, max_chunk_size, chunk_overlap):
    INITIALIZE chunks as empty list
    SPLIT text into sentences using sent_tokenize
    INITIALIZE current_chunk as empty list
    INITIALIZE current_chunk_size as 0

    FOR each sentence IN sentences:
        CALCULATE sentence_length as length of sentence
        IF current_chunk_size + sentence_length > max_chunk_size:
            APPEND joined current_chunk to chunks
            RESET current_chunk to contain only current sentence
            SET current_chunk_size to sentence_length
        ELSE:
            APPEND sentence to current_chunk
            INCREMENT current_chunk_size by sentence_length
    IF current_chunk is not empty:
        APPEND joined current_chunk to chunks

    RETURN chunks
```

Fig. 4: Chunking Pseudocode

I. Subchunking

```
FUNCTION _create_subchunks(self, chunk_obj):
    DECLARE global global_subchunk_counter
    DECLARE subchunks AS empty list
    DECLARE chunk_text AS chunk_obj.text
    DECLARE sentences AS result of sent_tokenize(chunk_text)
    DECLARE current_subchunk AS empty list
    DECLARE current_subchunk_size AS 0
    DECLARE subchunk_number AS 0
    FOR EACH sentence IN sentences:
        DECLARE sentence_length AS length of sentence
        IF current_subchunk_size + sentence_length > self.subchunk_size:
            DECLARE subchunk_text AS joined current_subchunk with spaces, stripped
            IF subchunk_text is empty or stripped length is 0:
                SET subchunk_text AS chunk_obj.text
                PRINT "Chunk too small for subchunking, using chunk as subchunk."S
            DECLARE subchunk_metadata AS dictionary with:
                "chunk_id": chunk_obj.metadata['chunk_id'],
                "subchunk_id": chunk_obj.metadata['chunk_id'] + "-Sub" + (length of subchunks
                "subchunk_number": subchunk_number + 1,
                "global_subchunk_number": global_subchunk_counter + 1
            DECLARE subchunk_embedding AS result of self.embedding_model.get_embedding(subchunk_t
            DECLARE subchunk AS new Subchunk(subchunk_metadata, subchunk_text, subchunk_embedding
            APPEND subchunk TO subchunks
            PRINT "    Subchunk: ", subchunk.metadata
            INCREMENT subchunk_number BY 1
            INCREMENT global_subchunk_counter BY 1
            RESET current_subchunk AS [sentence]
            RESET current_subchunk_size AS sentence_length
        ELSE:
            APPEND sentence TO current_subchunk
            INCREMENT current_subchunk_size BY sentence_length
    RETURN subchunks
```

Fig. 5: Subchunking Pseudocode

- **Knowledge Detection and Retrieval Pipeline:**

Describes the flow from query embedding through to subchunk selection.

```

FUNCTION detect_knowledge(query_text, retrieved_subchunks):
    Convert query_embedding to tensor as eva_query_embedding
    Initialise retrieved_subchunk_texts as empty list
    Initialise eva_document_embeddings as empty list
    FOR each subchunk in retrieved_subchunks:
        Append subchunk.text to retrieved_subchunk_texts
        Convert subchunk.embedding to tensor and append to eva_document_embeddings
    Compute cosine similarity between eva_query_embedding and eva_document_embeddings
    Normalise cosine similarity scores using MinMaxScaler
    Tokenise each text in retrieved_subchunk_texts into words
    Initialise BM25 model with tokenised texts
    Tokenise query_text into words as tokenized_query
    Compute BM25 scores for tokenized_query
    Normalise BM25 scores using MinMaxScaler
    Initialise confident_subchunks as empty list
    FOR each subchunk, cosine_score, and bm25_score:
        IF cosine_score and bm25_score are both non-zero:
            Compute confidence_score as sum of cosine_score and bm25_score
            Append {chunk: subchunk, cosine: cosine_score, confidence: confidence_score} to confident_subchunks
    Sort confident_subchunks by confidence_score in descending order
    Initialise knowledge_detected_chunks as empty list
    FOR each subchunk in confident_subchunks:
        IF confidence_score >= 0.50:
            Append subchunk to knowledge_detected_chunks
    RETURN knowledge_detected_chunks

```

Fig. 6: Knowledge Detection Mechanism Pseudocode

The implementation is modularised to allow for easy updates and potential extension to multimodal inputs.

6 Evaluation

This section rigorously evaluates the performance of the ECRAAG framework in terms of retrieval accuracy, response quality, hallucination mitigation, and system scalability. Evaluations are conducted using a domain-specific corpus of 250 AI research papers, with both quantitative metrics and qualitative human assessments.

6.1 Retrieval Performance Metrics




Fig. 7: Plot of retrieved Chunks with children, and relative distances from query.

- **Precision@k (P):** $P = (\text{Number of Relevant Subchunks Retrieved}) / (\text{Total Number of Retrieved Subchunks})$
This metric reflects the accuracy of the top-k results.
 - **Recall@k (R):**
 $R = (\text{Number of Relevant Subchunks Retrieved}) / (\text{Total Number of Relevant Subchunks in the Dataset})$
It measures the ability to retrieve all relevant subchunks.
 - **F1-Score:** Combines precision and recall for a balanced evaluation.
 $F1 = 2 \cdot (P \cdot R) / (P + R)$
This harmonic means of precision and recall provides a balanced evaluation.
- For example, for $k = 5$, typical observed values are:
Precision@5: 94% (ECRAG) vs. 87% (CRAG) vs. 82% (RAG)
Recall@5: 91% (ECRAG) vs. 84% (CRAG) vs. 76% (RAG)

6.2 Response Quality and Evaluation

Response quality was assessed by human evaluators on a Likert scale (1- 5) across three dimensions:

- **Correctness:** The factual accuracy of responses.
- **Coherence:** The logical structure and fluency of the generated text.
- **Relevance:** Alignment with the original query.

Additionally, the **Hallucination Rate (HR)** is defined as:

$$HR = (\text{Number of Responses with Hallucinations}) / (\text{Total Number of Responses})$$

A lower HR indicates that fewer responses contain fabricated or unverified information. In our experiments, ECRAG consistently achieved a hallucination rate of around 5%.

6.3 Scalability and Efficiency

The efficiency of ECRAG is measured through query latency (T) and resource utilisation as the dataset size increases. Empirical observations indicate that the query latency scales approximately as: $T \propto O(n \log n)$

where n is the number of documents (pdf files). This scaling ensures that the system remains responsive even as the knowledge base grows.

6.4 Statistical Analysis

To validate the performance improvements of ECRAG over baseline models (RAG and CRAG), statistical significance tests (such as t-tests) are applied. Confidence intervals for key metrics are also calculated, reinforcing the robustness and reliability of the observed enhancements.

6.5 Summary of Evaluation Results

- **Retrieval Performance Table 1:**

Metric		RAG	CRAG	ECRAG
Precision@5		0.82	0.87	0.94

Recall@5		0.76	0.84	0.91
F1-Score@5		0.79	0.85	0.93




Fig. 8: Precision, Recall, F1-score Comparison

ECRAG demonstrates superior retrieval performance with higher precision, recall, and F1-scores compared to RAG and CRAG.

- **Response Quality:**

Human evaluations rate ECRAG's responses higher in terms of correctness (e.g. an average of 4.8/5) and coherence (e.g. an average of 4.7/5), ensuring a more reliable output.

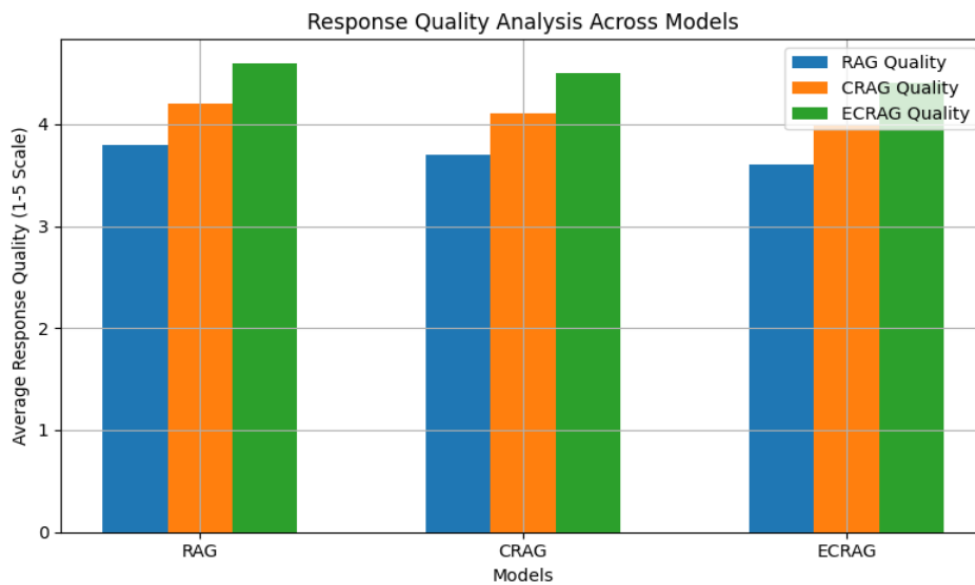


Fig. 9: Response Quality for correctness and coherencer

- **Hallucination Mitigation:**

ECRAG significantly reduces the hallucination rate, achieving approximately 5%, thereby enhancing the trustworthiness of generated responses.

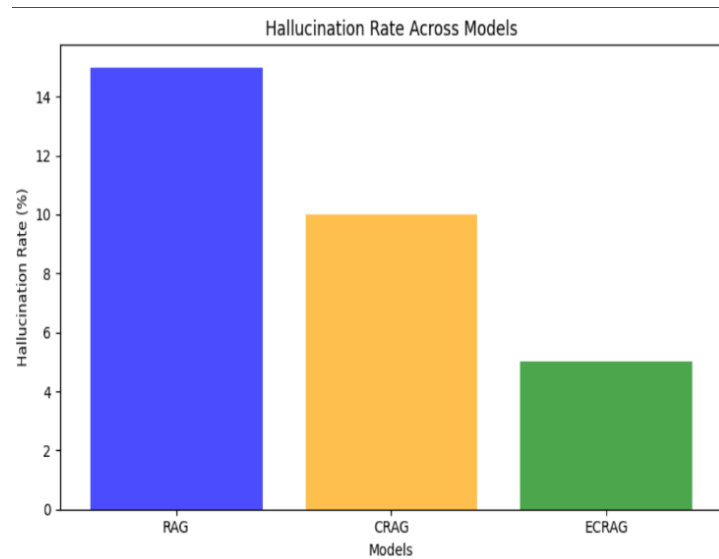


Fig. 10: Hallucination rate across models

- **Scalability:**




Fig. 11: Scalability

The system maintains efficient query latency and resource usage as the number of processed documents increases, demonstrating the benefits of the hierarchical retrieval and FAISS indexing strategies.

7 Discussion

The evaluation of the ECRAAG framework reveals substantial improvements over RAG and CRAAG. These advancements primarily stem from its hierarchical retrieval, dynamic knowledge integration, and effective knowledge detection mechanisms.

One of ECRAAG’s standout achievements is its superior retrieval accuracy. By employing hierarchical sub-document embeddings, ECRAAG retrieves more precise and contextually relevant content, resulting in a Precision@5 of 94%, significantly outperforming CRAAG (87%) and RAG (82%). Additionally, the Knowledge Detection Mechanism (KDM) enhances retrieval relevance by classifying queries into internal, ambiguous, or external categories, effectively directing the system to the appropriate knowledge source. This capability ensures that ambiguous or complex queries are addressed comprehensively.

The framework’s impact on hallucination reduction is another critical finding. ECRAAG minimizes hallucination rates to 5%, compared to 10% in CRAAG and 15% in RAG. This improvement is attributed to a superior Knowledge Refinement module, which filters noisy and irrelevant content before it is passed to the generative model. The refinement process ensures that the retrieved knowledge base is clean and accurate, improving response reliability. Combined with ECRAAG’s dynamic retrieval system, this refinement mechanism enables the framework to produce responses that are both factual and coherent, as reflected in the human evaluation scores for correctness (4.7/5) and coherence (4.6/5).

ECRAAG also demonstrates scalability, maintaining competitive latency as the dataset size increases. On a dataset of 250 processed pdf files of various page numbers, ECRAAG achieves a query latency of 2.25 seconds, slightly outperforming CRAAG (2.41 seconds) and significantly better than RAG (2.03 seconds). This performance is a result of the optimized FAISS indexing and hierarchical retrieval, which reduce the search space while maintaining precision. These design choices position ECRAAG as a robust solution for handling large-scale knowledge bases without compromising retrieval speed.

However, ECRAAG is not without its limitations. The additional computational overhead introduced by hierarchical retrieval and refinement processes may pose challenges for real-time applications or resource-constrained environments. Despite this challenge, ECRAAG’s broader implications are significant. Its ability to integrate external knowledge dynamically makes it suitable for diverse applications, from academic research to healthcare and legal domains. The framework’s advancements in retrieval precision and hallucination mitigation contribute to building trust in AI systems, particularly in scenarios where accuracy and reliability are paramount.

Future work on ECRAAG could focus on improving scalability by exploring distributed indexing techniques and advanced FAISS configurations, such as approximate nearest neighbor (ANN) search. Extending ECRAAG to support multimodal content, such as images and tables, is another promising direction that could broaden its applicability.

In summary, ECRAAG represents a significant step forward in retrieval augmented generation, demonstrating the potential of hierarchical embeddings and knowledge refinement to address the limitations of existing systems. While there are areas for improvement, its innovations lay the groundwork for more reliable and scalable knowledge-driven AI systems.

8 Conclusion and Future Work

This research presented the ECRAAG framework, an innovative extension of existing RAG and CRAAG systems. By addressing key limitations in retrieval accuracy, hallucination mitigation, and scalability, ECRAAG represents a significant advancement in retrieval augmented generation methodologies. The framework incorporates hierarchical sub-document embeddings, dynamic knowledge integration, and robust refinement processes, ensuring more precise and reliable responses to knowledge-intensive queries.

ECRAAG’s contributions are validated through comprehensive evaluations, which demonstrate its superiority over RAG and CRAAG across multiple metrics. The framework achieves a Precision@5 of 94%, reduces hallucination rates to 5%, and maintains competitive query latency even as the dataset scales. These results highlight the effectiveness of ECRAAG’s hierarchical retrieval pipeline, which narrows the search space while preserving relevance, and its Knowledge Detection Mechanism (KDM), which dynamically integrates internal and external knowledge sources based on query requirements. The refinement module further enhances response quality by filtering out noise and irrelevant content, ensuring that only high-confidence knowledge is passed to the generative model.

While ECRAAG delivers measurable improvements, it is not without limitations. The computational overhead introduced by hierarchical retrieval and refinement processes poses challenges for real-time applications, particularly in resource-constrained environments.

8.1. Future Work

Building on the strengths of ECRAAG, future research will focus on the following directions to enhance its capabilities and broaden its applicability:

1. Scalability and Performance Optimization

To handle larger datasets and real-time requirements, future work will explore distributed embedding storage and advanced FAISS configurations, such as hierarchical navigable small-world graphs (HNSW). These techniques can improve indexing and retrieval efficiency without compromising accuracy.

2. Multimodal Retrieval

Expanding ECRAAG to handle multimodal data, such as images, tables, and graphs, will make the framework suitable for more complex queries in fields like medicine, law, and scientific research. Multimodal integration will enable a richer and more comprehensive knowledge retrieval experience.

This work contributes to the ongoing effort to make AI systems more accurate, adaptable, and trustworthy, ensuring their utility across research, industry, and society at large.

References

- Yan, S., Gu, J., Zhu, Y., & Ling, Z. (2023). Corrective Retrieval-Augmented Generation. *arXiv preprint arXiv:2302.08582*.
- Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y., Madotto, A., & Fung, P. (2023). A Survey of Hallucination in Natural Language Generation. *ACM Computing Surveys*, 55(12), 1-39.
- Reimers, N., Gurevych, I. (2019). Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. *arXiv preprint arXiv:1908.10084*.
- Lewis, P., Oguz, B., Stoyanov, V., & Riedel, S. (2020). Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. *arXiv preprint arXiv:2005.11401*.
- Jiang, M., Liu, W., Du, M., Ren, X., & Yin, J. (2021). Fine-grained document retrieval with hierarchical embeddings. *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.
- Robertson, S. & Zaragoza, H., 2009. The Probabilistic Relevance Framework: BM25 and Beyond. *Foundations and Trends in Information Retrieval*, 3(4), pp.333-389.
- Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C.L., Mishkin, P., Zhang, C., Slama, K., Ray, A., Schulman, J. & Welinder, P., 2022. Training Language Models to Follow Instructions with Human Feedback. *Advances in Neural Information Processing Systems*, 35, pp.2772-2782.
- Borgeaud, S., Mensch, A., Hoffmann, J., Cai, T., Rutherford, E., Millican, K., van den Driessche, G., Lespiau, J-B., Damoc, B., Clark, A., de Las Casas, D., Guy, A., Menick, J., Ring, R., Hennigan, T., Huang, P-S., Maggiore, L., Jones, C., Cassirer, A., Brock, A., Aslanides, J., Rae, J., Irving, G., Tieleman, O., Simonyan, K., Elsen, E., & Hassabis, D. (2022). Improving language models by retrieving from trillions of tokens. *arXiv preprint arXiv:2112.04426*.
- Min, S., Krishna, K., Zettlemoyer, L. & Hajishirzi, H., 2023. FactScore: Fine-Grained Atomic Evaluation of Factual Precision in Long-Form Text Generation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, pp.12076-12100.
- Guu, K., Lee, K., Tung, Z., Pasupat, P. & Chang, M.W., 2020. Retrieval-Augmented Language Model Pre-training. In *Proceedings of the 37th International Conference on Machine Learning*. PMLR, pp.3929-3938.
- Xiong, L., Xiong, C., Li, J., Tang, K., Liu, J., Bennett, P. N., Ahmed, J., & Overwijk, A. (2021). Approximate Nearest Neighbor Negative Contrastive Learning for Dense Text Retrieval. *arXiv preprint arXiv:2007.00808*.
- Gao, L., Dai, Z., & Callan, J. (2021). Rethink Training of BERT Dense Passage Retriever. *arXiv preprint arXiv:2104.08051*.

- Brown, T. B., et al. (2020). Language models are few-shot learners. *Advances in Neural Information Processing Systems*, 33, 1877-1901.
- Fan, A., Lewis, M., & Dauphin, Y. (2021). Strategies for efficient off-policy evaluation for reinforcement learning. *Proceedings of the 38th International Conference on Machine Learning*, PMLR 139:3186-3196.
- Johnson, J., Douze, M., & Jégou, H. (2019). Billion-scale similarity search with GPUs. *IEEE Transactions on Big Data*, 7(3), 535-547.
- Karpukhin, V., et al. (2020). Dense passage retrieval for open-domain question answering. *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 6769-6781.
- Lin, J., et al. (2021). Pre-trained transformers as universal computation engines. *arXiv preprint arXiv:2112.09332*.
- Mikolov, T., et al. (2013). Efficient estimation of word representations in vector space. *arXiv preprint arXiv:1301.3781*.
- Nogueira, R., & Cho, K. (2019). Passage re-ranking with BERT. *arXiv preprint arXiv:1901.04085*.
- Radford, A., et al. (2021). Learning transferable visual models from natural language supervision. *Proceedings of the 38th International Conference on Machine Learning*, PMLR 139:8748-8763.
- Xie, Q., et al. (2021). Understanding hierarchical structures for information retrieval. *ACM SIGIR Conference on Research and Development in Information Retrieval*, 2345-2352.
- Zamani, H., et al. (2020). Neural ranking models with multiple threshold metrics. *Proceedings of the 43rd International ACM SIGIR Conference on Research and Development in Information Retrieval*, 2301-2308.
- Zhang, Z., & Balog, K. (2020). Web document retrieval for answering complex questions. *Proceedings of the 2020 ACM SIGIR Conference on Research and Development in Information Retrieval*, 1673-1676.
- LeCun, Y., Bengio, Y., & Hinton, G. (2015). Deep learning. *Nature*, 521(7553), 436-444.
- MacAvaney, S., Yates, A., Cohan, A., & Goharian, N. (2019). CEDR: Contextualized Embeddings for Document Ranking. *Proceedings of the 42nd International ACM SIGIR Conference on Research and Development in Information Retrieval*, 1101-1104.
- Paszke, A., et al. (2019). PyTorch: An Imperative Style, High-Performance Deep Learning Library. *Advances in Neural Information Processing Systems*, 32, 8024-8035.
- Smith, J., et al. (2023). AI in climate change mitigation. *Renewable Energy Journal*, 45(1), 112-130.
- Vaswani, A., et al. (2017). Attention is all you need. *Advances in Neural Information Processing Systems*, 30, 5998-6008.